

A Split-and-Recombine Approach for Follow-up Query Analysis

Speaker: Qian Liu (刘乾)

Date : 2019.12.22

Outline

1. Introduction
2. Methods
3. Our Research
4. Take-home-message
5. Reference

1. Introduction

Database

Brand	Sales	Profit	Year
BMW	31020	5000	2009
Ford	25220	3000	2009
Benz	47060	6000	2009

Text-to-SQL



show the sales of BMW in 2009

SELECT Sales WHERE Brand = BMW and Year = 2009



Conversation



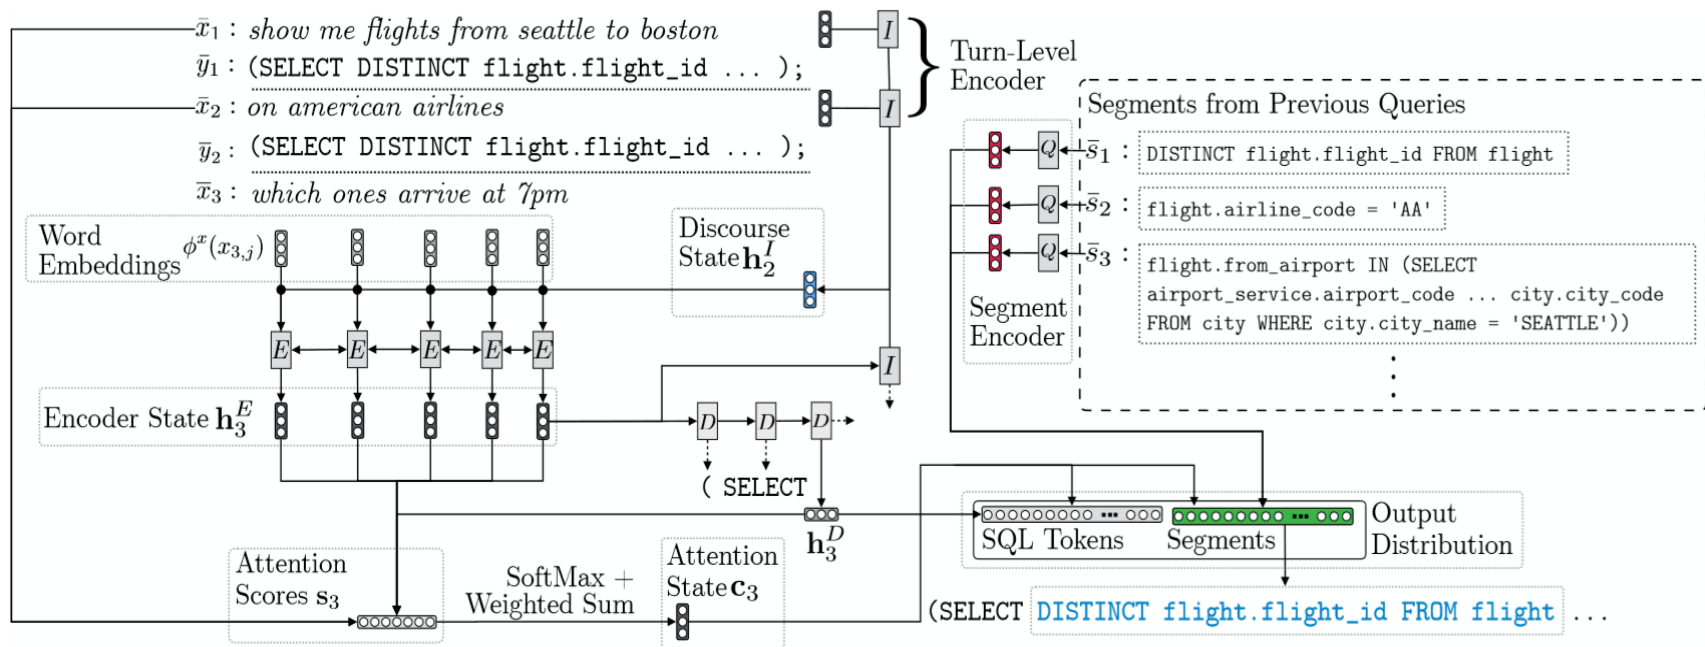
what about profit ?

SELECT Profit WHERE Brand = BMW and Year = 2009



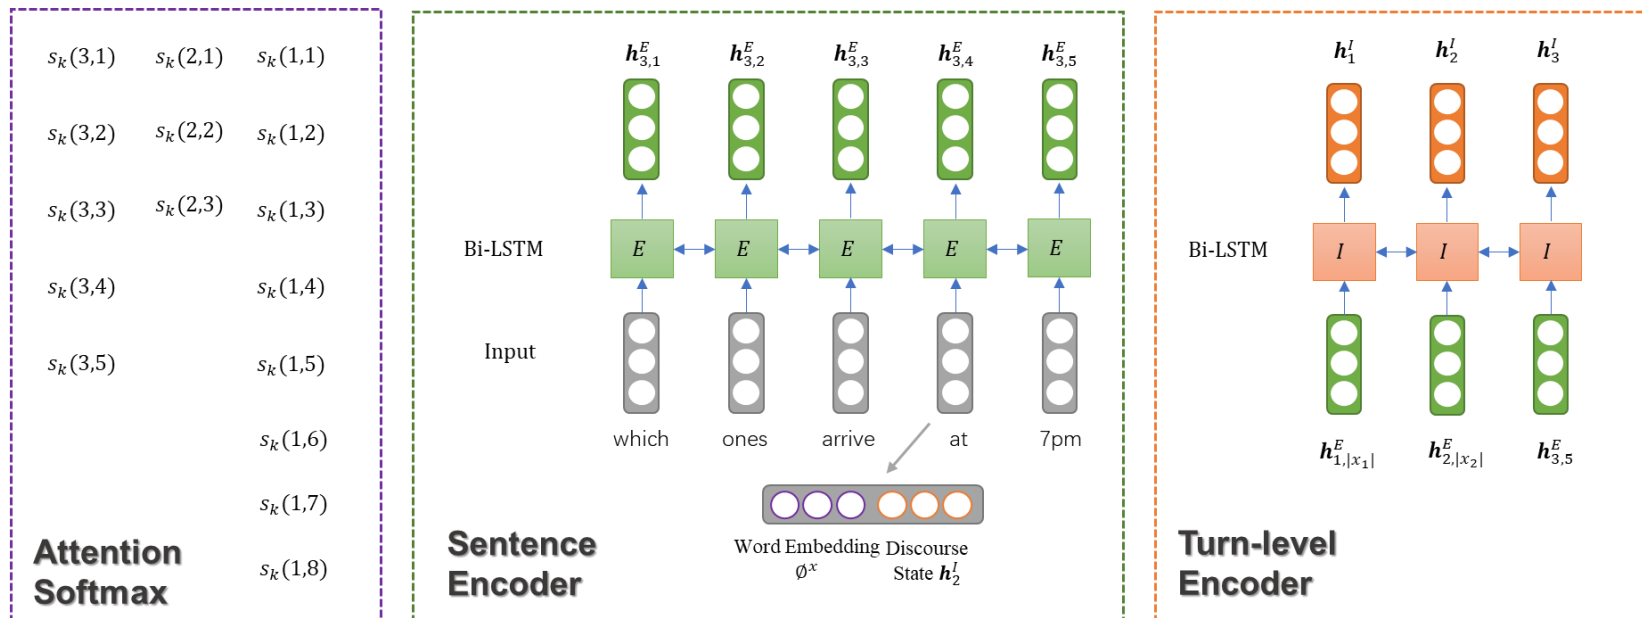
2. Methods

Traditional methods always model the problem as an end-to-end manner^[1].



2. Methods

- Turn-level encoder:** hierarchically represent the dialogue history



2. Methods

- **SQL Segment Encoder:** reduce decoding steps by a large margin

```
(SELECT DISTINCT flight.flight id FROM flight WHERE
(flight.airline code = 'AA') AND (flight.from
airport IN (SELECT airport service.airport code FROM
airport service WHERE airport service.city code IN
(SELECT city.city code FROM city WHERE city.city
name = 'SEATTLE')))) AND (flight.to airport IN
(SELECT air port service.airport code FROM airport
service WHERE airport service.city code IN (SELECT
city.city code FROM city WHERE city.city name
= 'BOSTON')))) AND (flight.flight days IN (SELECT
days.days code FROM days WHERE days.day name IN
(SELECT date day.day name FROM date day WHERE date
day.year = 1993 AND date day.month number = 2 AND
date day.day number = 8))));
```

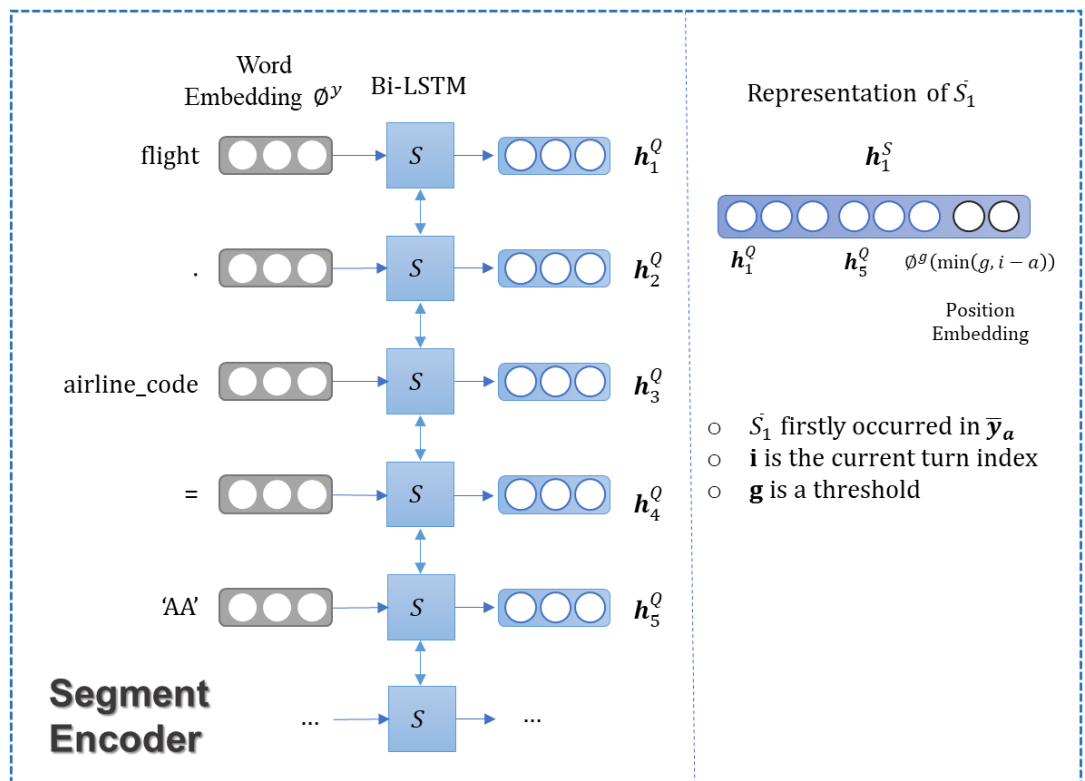
Construct Tree &
Extract Subtrees

flight.airline_code = 'AA'

DISTINCT flight.flight_id FROM
flight

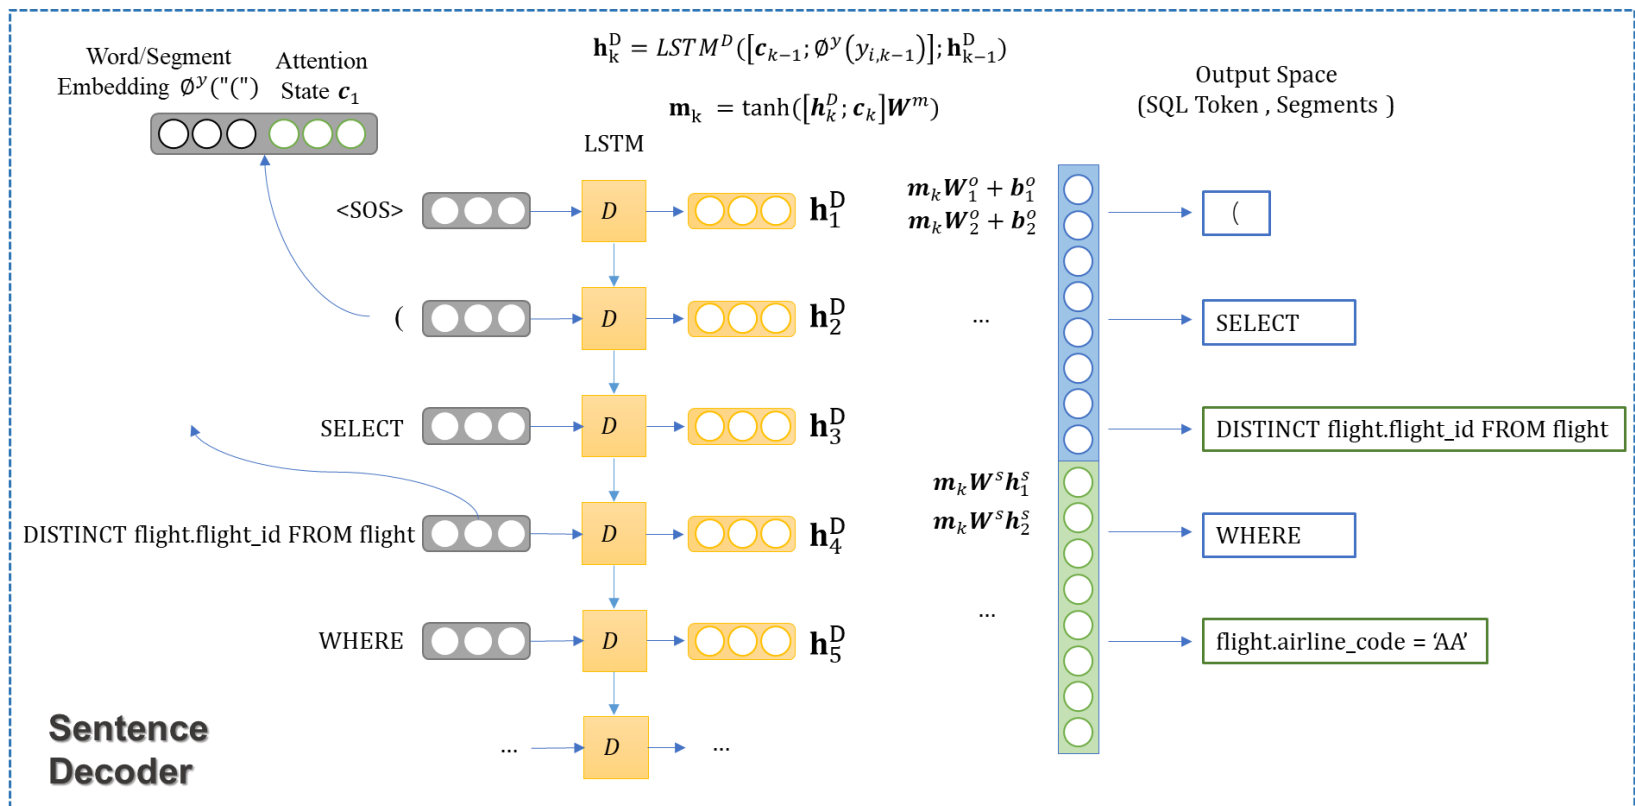
flight.from_airport IN (... WHERE
city.city_name = 'SEATTLE'))

...



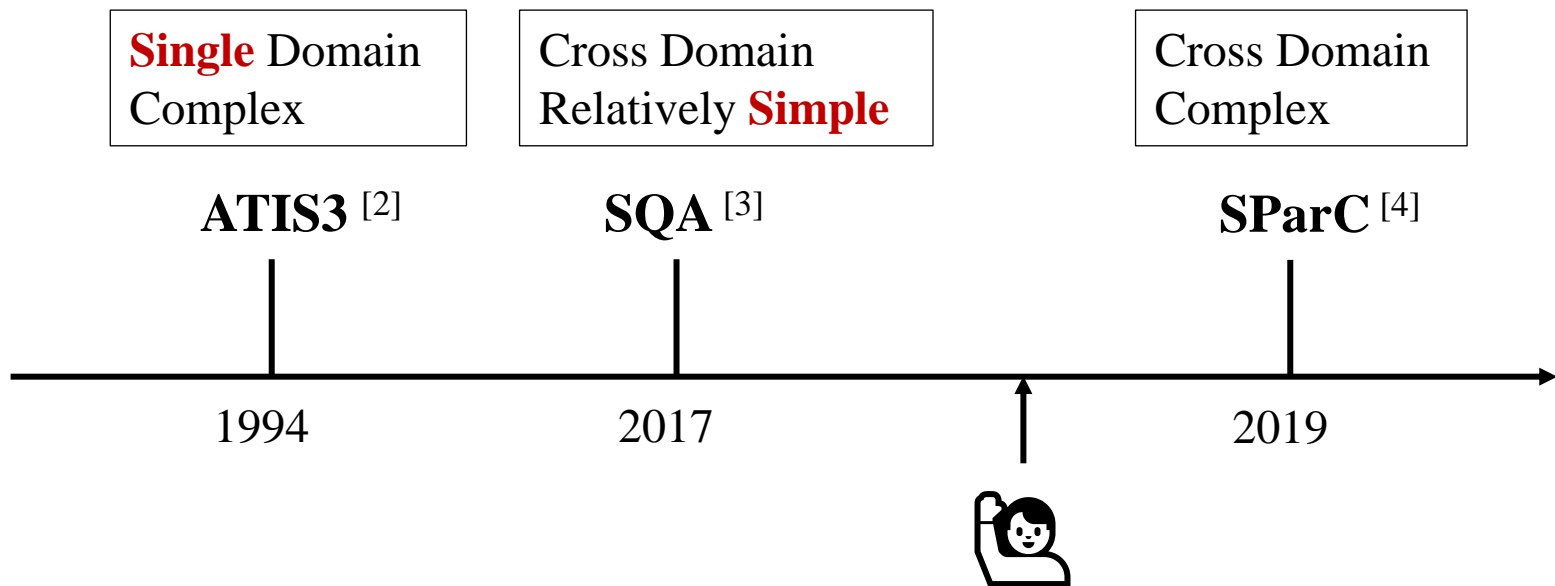
2. Methods

- Joint Decoding:** decode SQL Token & Segment jointly



2. Methods

Previous works focus on end-to-end context-dependent parsing (e.g. Conversational Text-to-SQL). However, related datasets are very rare 🙄.



2. Methods

However, semantic parsing techniques obtained a large attention.
Therefore, we have a question:

Could we use trained semantic parsers to solve
the context-dependent semantic parsing problem?

2. Methods

However, semantic parsing techniques obtained a large attention.
Therefore, we have a question:

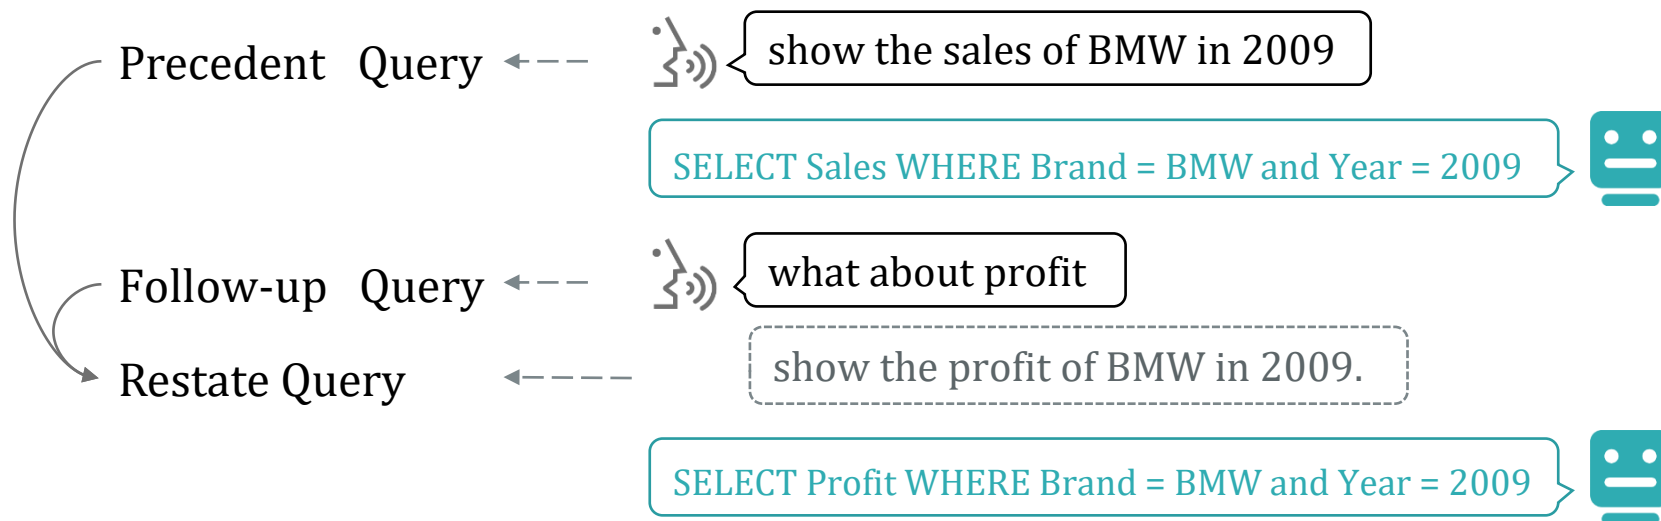
Could we use trained semantic parsers to solve
the context-dependent semantic parsing problem?

Yes, we could make it when imposing an auxiliary task

Follow-up Query Analysis (FQA)

2. Methods

Brand	Sales	Profit	Year
BMW	31020	5000	2009
Ford	25220	3000	2009
Benz	47060	6000	2009



2. Methods

- Previous works aims to achieve the task end-to-end.
- We aim to use a pretrained semantic parser to do it.
- The annotation of FQA is cheaper than conversational semantic parsing.

Previous Work

Context-dependent
Parsing

Our Method

Follow-up Query
Analysis



Parsing

2. Methods

Scenario	Example
Analytics	Precedent : In 1995, is there any network named CBC?
	Follow-up : Any TSN?
	Fused : In 1995, is there any network named TSN?
Compare	Precedent : How much money has Smith earned?
	Follow-up : Compare it with Bill Collins.
	Fused : Compare money Smith earned with Bill Collins.
Calc & Stats	Precedent : List all universities founded before 1855.
	Follow-up : Show their number.
	Fused : Show the number of all universities founded before 1855.
Extremum	Precedent : Which stadium has the most capacity?
	Follow-up : Which get the highest attendance?
	Fused : Which stadium get the highest attendance?
Filter	Precedent : How many roles are from studio paramount?
	Follow-up : List all titles produced by that studio.
	Fused : List all titles produced by studio paramount.
Group	Precedent : Show the industry which has the most companies?
	Follow-up : Show in different countries.
	Fused : Show the industry which has the most companies in different countries.
Sort	Precedent : Show all chassis produced after the year 1990.
	Follow-up : Sort them by year.
	Fused : Show all chassis produced after the year 1990 and sort by year.
Search	Precedent : What position did Sid O'Neill play?
	Follow-up : Which players else are in the same position?
	Fused : Which players play in the position of Sid O'Neill excluding Sid O'Neill?

2. Methods

The challenge turns to be:
How to perform the task FQA ?

2. Methods

The restate query always has **a large overlap** with the input queries, hence we have following options:

- Sequence to sequence with attention
- Copy mechanism (CopyNet/Point Generator) [5]
- Concatenate

show the sales of BMW in 2009

what about profit

} → show the profit of BMW in 2009

2. Methods

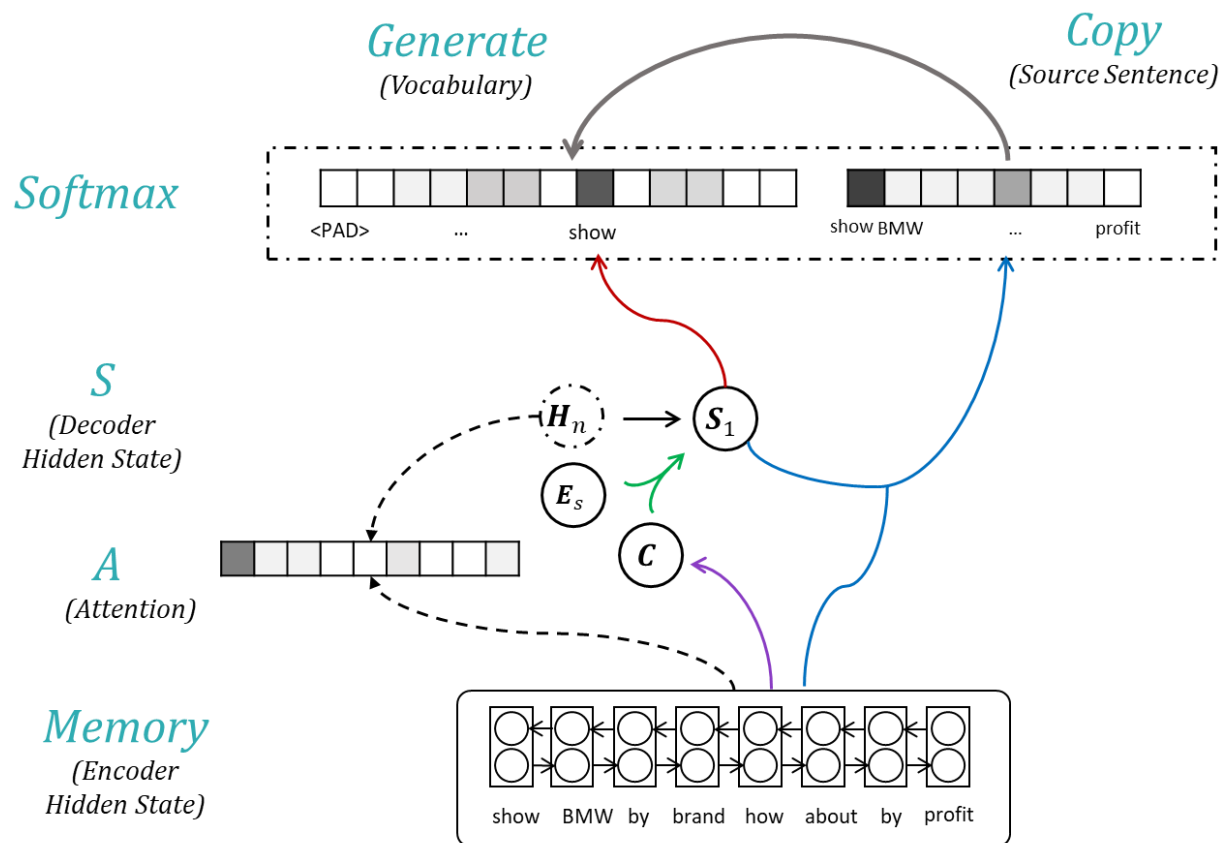
The restate query contains **anaphora** which refers to the entities in the input queries, we consider:

- End to end coreference resolution [6]

How much money has Bill earned
Compare it with Smith } → Compare money Bill has earned with Smith

2. Methods

COPYNET



$$\alpha_i = V_a^T \tanh(W[H_n H_i])$$

$$a_i = \frac{e^{\alpha_i}}{\sum e^{\alpha_k}}$$

$$C = \sum a_i \cdot H_i$$

$$S_1 = RNN(H_n, [E_s, C])$$

$$O_G = W_o S_1$$

$$\beta_j = \tanh(W_c H_j) S_1$$

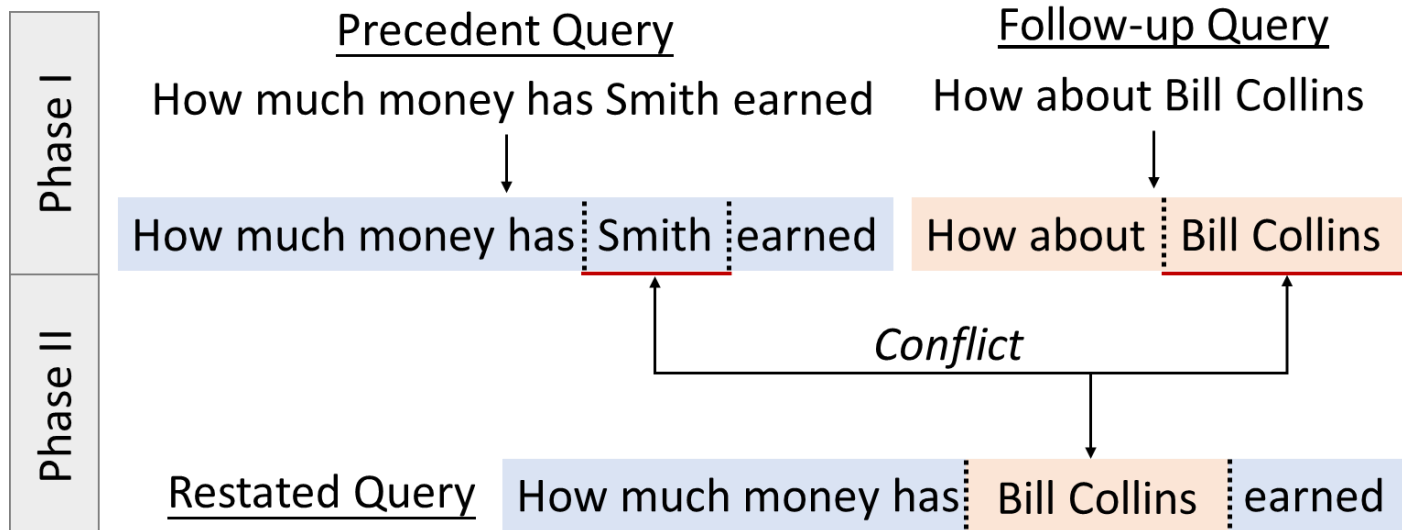
$$O_C = [\beta_1, \dots, \beta_s]$$

$$P = \text{Softmax}([O_G, O_C])$$

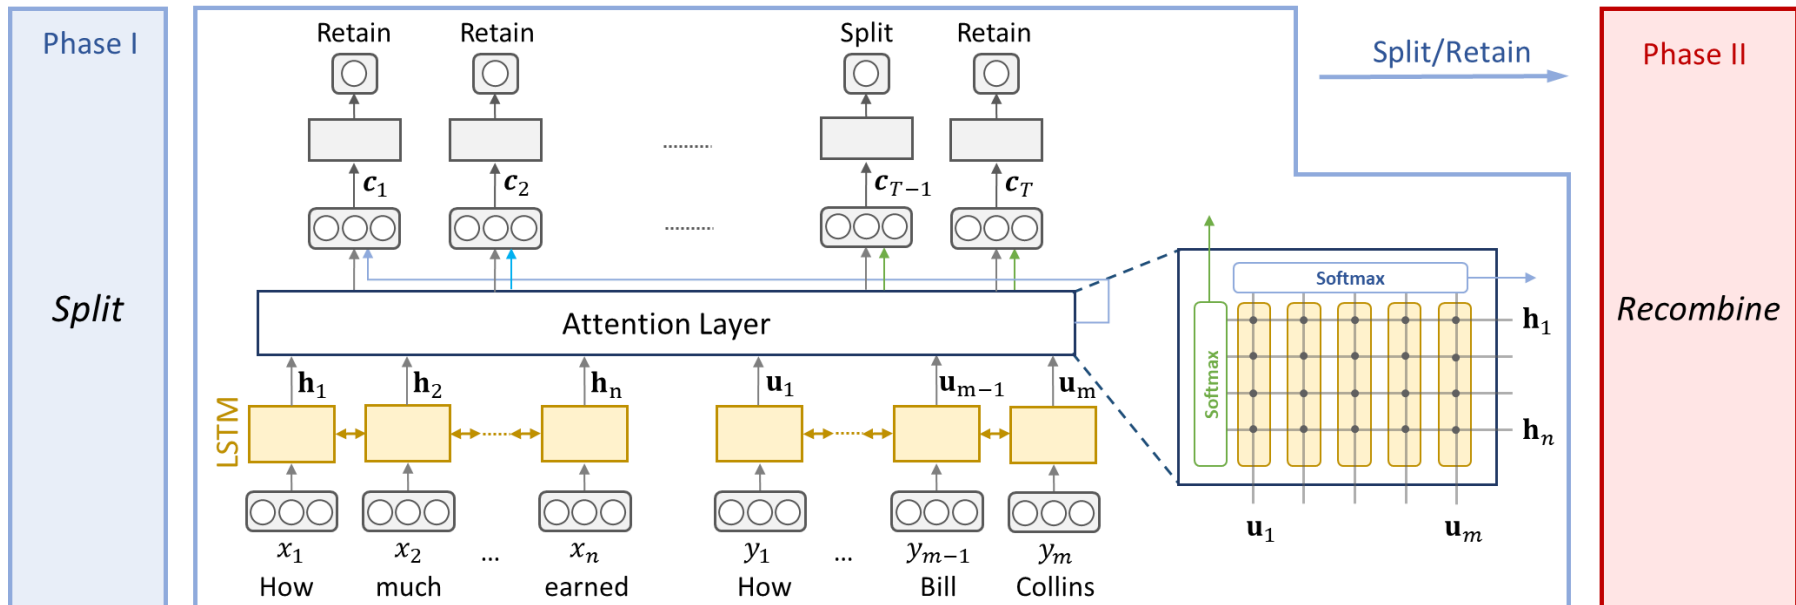
$$p(y_t) = p(y_t, g | \cdot) + p(y_t, c | \cdot)$$

3. Ours Research

Our solution: ***Split*** and ***Recombine***



3. Ours Research



- ***Split*** is modeled as a sequence labeling problem $\{0,1\}$.
- ***Recombine*** generates the restate based on semantic conflicting.

3. Ours Research


The architecture of our **SplitNet** is as follows:

- Embedding Layer: character, word and segment $[\phi_c; \phi_w; \phi_s]$
- Context Layer: Bi-LSTM $\vec{h}_i = \overrightarrow{\text{LSTM}}(\phi(x_i); \vec{h}_{i-1})$

3. Ours Research

The architecture of our **SplitNet** is as follows:

- Embedding Layer: character, word and segment $[\phi_c; \phi_w; \phi_s]$
- Context Layer: Bi-LSTM $\vec{h}_i = \overrightarrow{\text{LSTM}}(\phi(x_i); \vec{h}_{i-1})$
- Bi-Attention Layer: precedent-to-follow attention & follow-to-precedent attention



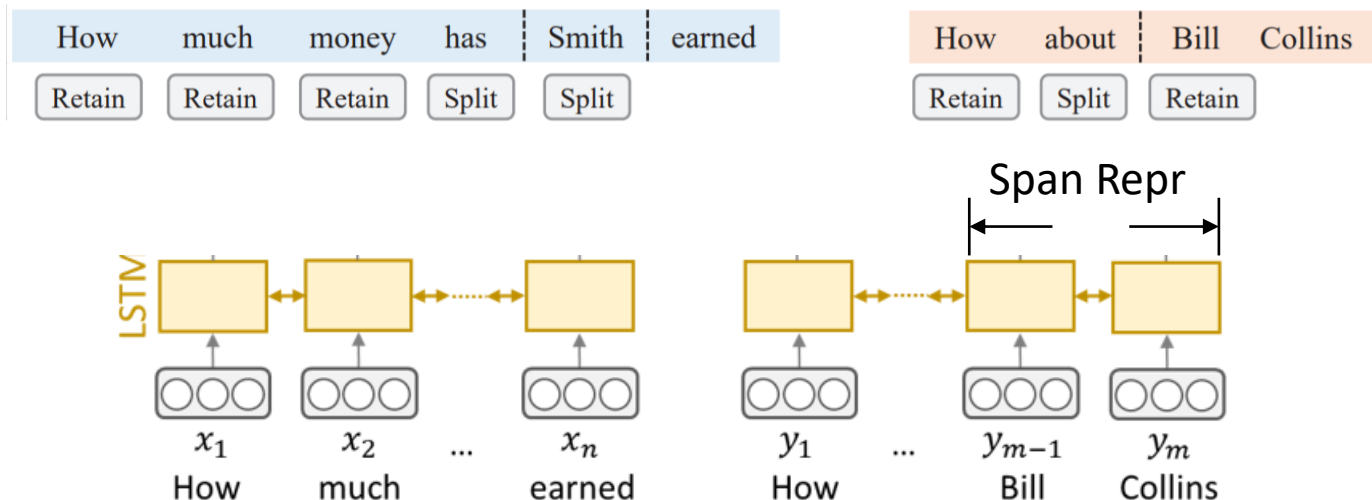
- Output: combine adjacent states to predict the action

$$\mathbf{c}^{x_i} = [\mathbf{h}_i; \mathbf{h}_i \circ \tilde{\mathbf{h}}_i; \mathbf{h}_{i+1} \circ \tilde{\mathbf{h}}_{i+1}] \quad \sigma(\mathbf{W} * \mathbf{c}_t + b)$$

3. Ours Research

The recombination process is to generate the restate query:

1. Giving spans, we could obtain the representations of them.

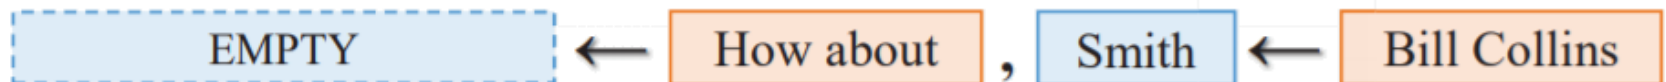


3. Ours Research

- Using the representation, we could identify whether the spans semantically conflict.

	How much money has	Smith	earned
How about			
Bill Collins			

- Then use the 1-1 replacement principle, we could get the result.



3. Ours Research

An ideal solution, but there is no annotation!

- **We do not know** how to split the utterance to spans.
- **We do not know** which span conflicts with which one.

3. Ours Research

An ideal solution, but there is no annotation!

- **We do not know** how to split the utterance to spans.
- **We do not know** which span conflicts with which one.

We propose to solve the problem using RL.

- **Pretrain** the model under noisy distant supervision.
- **Fine-tune** it with rewards using policy gradient.

3. Ours Research

Find the **maximum common string** to export the training data

How much money has Smith earned

How about Bill Collins

How much money has Bill Collins earned

3. Ours Research

A simple way to fine-tune the model is to sample both ***Split*** and ***Recombine***. The reward could be BLEU or other task-related metrics.

$$\mathcal{L}_{rl} = \mathbb{E} \left[\sum_{\tilde{\mathbf{z}} \in \mathcal{Z}} \sum_{q \in \mathcal{Q}} P_{\text{split}}(q | \mathbf{x}, \mathbf{y}) P_{\text{rec}}(\tilde{\mathbf{z}} | q) r(\mathbf{z}, \tilde{\mathbf{z}}) \right]$$

3. Ours Research



A simple way to fine-tune the model is to sample both ***Split*** and ***Recombine***. The reward could be BLEU or other task-related metrics.

$$\mathcal{L}_{rl} = \mathbb{E} \left[\sum_{\tilde{\mathbf{z}} \in \mathcal{Z}} \sum_{q \in \mathcal{Q}} P_{\text{split}}(q | \mathbf{x}, \mathbf{y}) P_{\text{rec}}(\tilde{\mathbf{z}} | q) r(\mathbf{z}, \tilde{\mathbf{z}}) \right]$$

But sadly 🤔, it convergences **slow** and performs **bad**.

3. Ours Research

Why ?

- **The space Q:** $2^{|Sequence|}$ 
- **The space Z:** span-based permutation 

The space $Q \times Z$ is vast, it is hard to train well via direct sampling.

3. Ours Research

Is there any way to
reduce the sampling space?

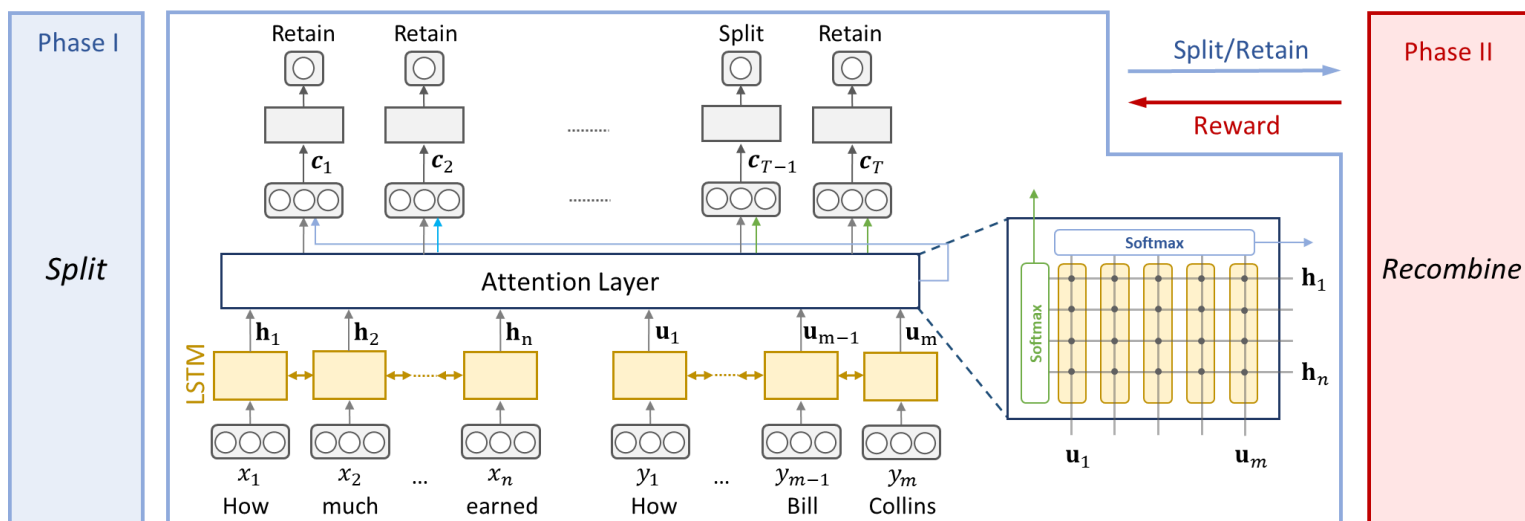
3. Ours Research

Yes, if we enumerate the space Z .

3. Ours Research

Yes, if we enumerate the space Z.

Reward shifting :Enumerating the space Z (a small effort) to provide rewards for action sampling in the space Q.



3. Ours Research

Recalling the equation of objective, the reward comes out:

$$\mathcal{L}_{rl} = \mathbb{E} \left[\sum_{\tilde{\mathbf{z}} \in \mathcal{Z}} \sum_{q \in \mathcal{Q}} P_{\text{split}}(q | \mathbf{x}, \mathbf{y}) P_{\text{rec}}(\tilde{\mathbf{z}} | q) r(\mathbf{z}, \tilde{\mathbf{z}}) \right],$$



$$\mathcal{L}_{rl} = \mathbb{E} \left[\sum_{q \in \mathcal{Q}} P_{\text{split}}(q | \mathbf{x}, \mathbf{y}) \sum_{\tilde{\mathbf{z}} \in \mathcal{Z}} P_{\text{rec}}(\tilde{\mathbf{z}} | q) r(\mathbf{z}, \tilde{\mathbf{z}}) \right],$$

3. Ours Research

Recalling the equation of objective, the reward comes out:

$$\mathcal{L}_{rl} = \mathbb{E} \left[\sum_{\tilde{\mathbf{z}} \in \mathcal{Z}} \sum_{q \in \mathcal{Q}} P_{\text{split}}(q | \mathbf{x}, \mathbf{y}) P_{\text{rec}}(\tilde{\mathbf{z}} | q) r(\mathbf{z}, \tilde{\mathbf{z}}) \right],$$

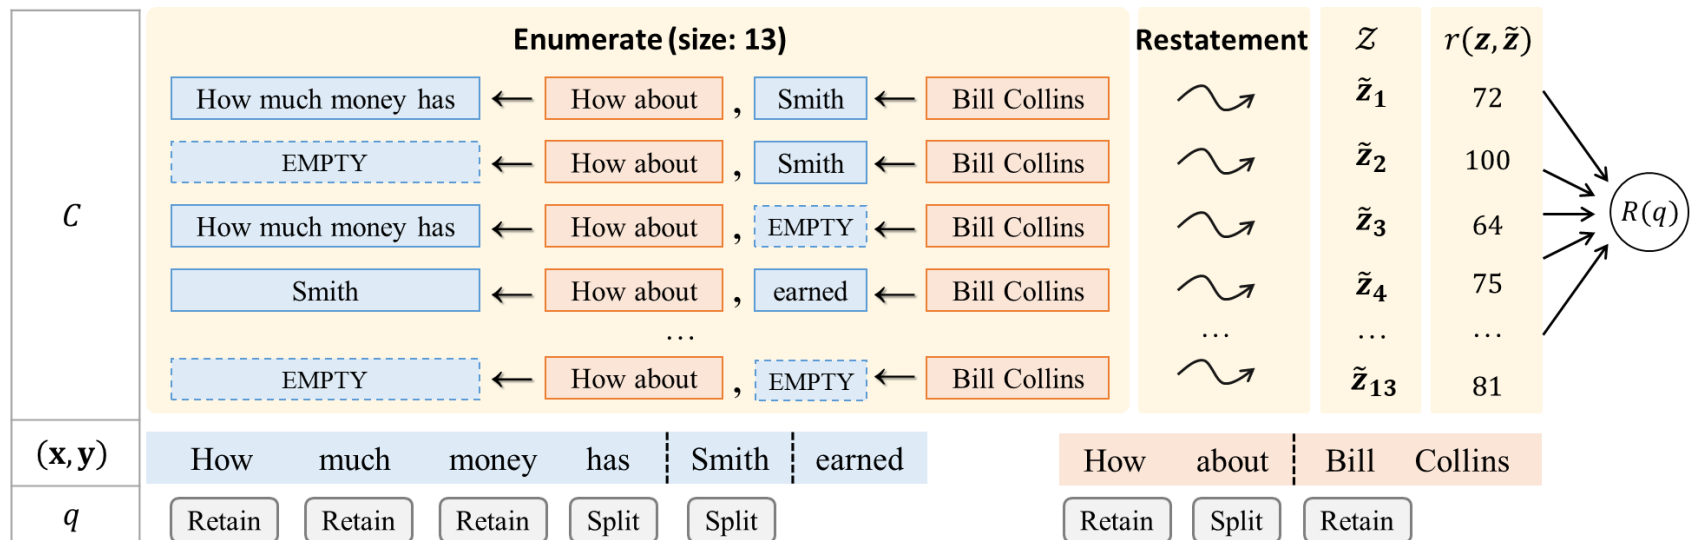


$$\mathcal{L}_{rl} = \mathbb{E} \left[\sum_{q \in \mathcal{Q}} P_{\text{split}}(q | \mathbf{x}, \mathbf{y}) \sum_{\tilde{\mathbf{z}} \in \mathcal{Z}} P_{\text{rec}}(\tilde{\mathbf{z}} | q) r(\mathbf{z}, \tilde{\mathbf{z}}) \right]$$

Reward for **SplitNet** $R(q, \mathbf{z})$

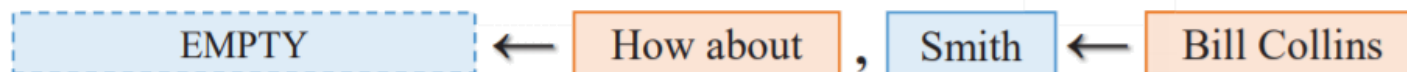
3. Ours Research

- Specifically, we systemically enumerate all valid **candidates**, and obtain the **R**.
- For unexpected large Z , we directly return the reward 0.



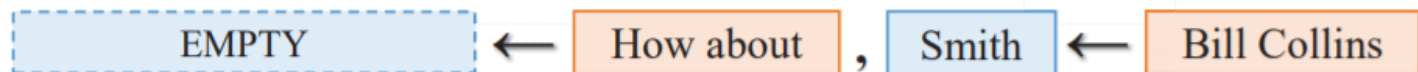
3. Ours Research

- The next question is, how to train the ***recombine***, a neural process where the span representations are learned.
- We take the candidate which gains the highest reward and regard it as the ground-truth.



3. Ours Research

- The next question is, how to train the ***recombine***, a neural process where the span representations are learned.
- We take the candidate which gains the highest reward and regard it as the ground-truth.



Objective Function

$$\cos[\text{Smith}, \text{Bill Collins}] > \cos[\text{EMPTY}, \text{Bill Collins}]$$

3. Ours Research


We mainly conducted experiments on **FollowUp** dataset [7]. Following metrics are employed to evaluate our method:

- ✓ **SymAcc**: detects whether all the SQL-related words are correctly involved in the predicted queries.
- ✓ **BLEU**: evaluates how similar the predicted queries are to golden ones.
- ✓ **AnsAcc**: check the answer accuracy of predicted queries manually.


Model	Dev		Test		
	SymAcc (%)	BLEU (%)	SymAcc (%)	BLEU (%)	AnsAcc (%)
SEQ2SEQ [†] (Bahdanau et al., 2015)	0.63 \pm 0.00	21.34 \pm 1.14	0.50 \pm 0.22	20.72 \pm 1.31	–
COPYNET [†] (Gu et al., 2016)	17.50 \pm 0.87	43.36 \pm 0.54	19.30 \pm 0.93	43.34 \pm 0.45	–
COPY+BERT (Devlin et al., 2019)	18.63 \pm 0.61	45.14 \pm 0.68	22.00 \pm 0.45	44.87 \pm 0.52	–
CONCAT [†]	–	–	22.00 \pm –	52.02 \pm –	25.24
E2ECR [†] (Lee et al., 2017)	–	–	27.00 \pm –	52.47 \pm –	27.18
FANDA [†] (Liu et al., 2019)	49.00 \pm 1.28	60.14 \pm 0.98	47.80 \pm 1.14	59.02 \pm 0.54	60.19
STAR	55.38 \pm 1.21	67.62 \pm 0.65	54.00 \pm 1.09	67.05 \pm 1.05	65.05

3. Ours Research

We also do ablation study to identify the modules of our method.



3. Ours Research



3. Ours Research

Our model could be incorporated with any parser. Thus, we performed extensive experiments on the **SQA** dataset ^[3].

Model	Precedent	Follow-up
DynSP (Iyyer et al., 2017)	70.9	35.8
NP (Neelakantan et al., 2016)	58.9	35.9
NP + STAR	58.9	38.1
DynSP + STAR	70.9	39.5
DynSP* (Iyyer et al., 2017)	70.4	41.1

3. Ours Research

Here are three real cases. Blue spans are those that have conflicts, and histograms represent the conflict probabilities to all the spans in precedent queries.

No	Case Analysis	
1	<i>Precedent</i>	: [compared to glebe park] [, does] [hampden park] [holds more attendances at capacity ?]
	<i>Follow-up</i>	: [how about] [compared to balmoor ■■■]
	STAR	: compared to balmoor , does hampden park holds more attendances at capacity ?
2	<i>Precedent</i>	: [Is there any book which belongs to] [Nancy miller]
	<i>Follow-up</i>	: [I mean] [the writer Nancy miller ■■]
	STAR	: Is there any book which belongs to the writer Nancy miller
3	<i>Precedent</i>	: [show directors of] [greatest love and promised land]
	<i>Follow-up</i>	: [show air date of ■■] [those two films ■■]
	STAR	: show air date of greatest love and promised land

3. Ours Research

Our code and data are released in Github:

- <https://github.com/microsoft/EMNLP2019-Split-And-Recombine>
- <https://github.com/SivilTaram/FollowUp>

4. Take-home-message

- For conversational scenarios, we can combine pretrained context-independent models and the task of FQA to handle the scenarios.
- Imposing a latent intermediate structure encounters the problem of hard training. We may do reward shifting to avoid large sampling space with a small enumerating effort.

5. References

- [1] Suhr A, Iyer S, Artzi Y. Learning to Map Context-Dependent Sentences to Executable Formal Queries
- [2] Deborah A. Dahl, Madeleine Bates, Michael Brown, William M. Fisher, Kate Hunicke-Smith, David S. Pallett, Christine Pao, Alexander I. Rudnicky, and Elizabeth Shriberg. 1994. Expanding the scope of the ATIS task: The ATIS-3 corpus.
- [3] Mohit Iyyer, Wen-tau Yih, and Ming-Wei Chang. 2017. Search-based neural structured learning for sequential question answering.
- [4] Yu T, Zhang R, Yasunaga M, et al. SParC: Cross-Domain Semantic Parsing in Context.
- [5] Gu J, Lu Z, Li H, et al. Incorporating copying mechanism in sequence-to-sequence learning.

5. References

- [6] Kenton Lee, Luheng He, Mike Lewis, and Luke Zettlemoyer. 2017. End-to-end neural coreference resolution.
- [7] Qian Liu, Bei Chen, Jian-Guang Lou, Ge Jin, and Dongmei Zhang. 2019. FANDA: A novel approach to perform follow-up query analysis.

Thanks

Speaker : Qian Liu